

## ✓ Business Case: Delhivery - Feature Engineering

### DELIVERY DATA ANALYSIS:

This case study focuses on the logistics sector, specifically examining a leading Indian company in this domain. The study aims to enhance operational efficiency and service quality through in-depth data analysis.

### ✓ OBJECTIVE

The main objective of feature engineering for Delhivery is to identify, create, and optimize relevant features from the available data sources. This will help in improving the accuracy and effectiveness of various predictive models and analytical tools used in delivery operations. This involves cleaning, manipulating, and understanding large sets of logistics data.

```
from google.colab import files
uploaded = files.upload()

 No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving delhivery_data.csv to delhivery_data.csv

import pandas as pd
df = pd.read_csv("delhivery_data.csv")
df.head()
```

	data	trip_creation_time	route_schedule_uuid	route_type	trip_uuid	source_center	source_name	destin
0	training	35:36.5	thanos::sroute:eb7bfc78-b351-4c0e-a951-fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)	I
1	training	35:36.5	thanos::sroute:eb7bfc78-b351-4c0e-a951-fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)	I
2	training	35:36.5	thanos::sroute:eb7bfc78-b351-4c0e-a951-fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)	I
3	training	35:36.5	thanos::sroute:eb7bfc78-b351-4c0e-a951-fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)	I
4	training	35:36.5	thanos::sroute:eb7bfc78-b351-4c0e-a951-fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)	I

5 rows × 24 columns

df.shape

(7178, 24)

df.columns

```
Index(['data', 'trip_creation_time', 'route_schedule_uuid', 'route_type',
       'trip_uuid', 'source_center', 'source_name', 'destination_center',
       'destination_name', 'od_start_time', 'od_end_time',
       'start_scan_to_end_scan', 'is_cutoff', 'cutoff_factor',
       'cutoff_timestamp', 'actual_distance_to_destination', 'actual_time',
       'osrm_time', 'osrm_distance', 'factor', 'segment_actual_time',
       'segment_osrm_time', 'segment_osrm_distance', 'segment_factor'],
      dtype='object')
```

df.dtypes

data	object
trip_creation_time	object
route_schedule_uuid	object
route_type	object
trip_uuid	object

```

source_center          object
source_name            object
destination_center     object
destination_name       object
od_start_time          object
od_end_time            object
start_scan_to_end_scan float64
is_cutoff              bool
cutoff_factor          int64
cutoff_timestamp       object
actual_distance_to_destination float64
actual_time             float64
osrm_time               float64
osrm_distance           float64
factor                 float64
segment_actual_time    float64
segment_osrm_time      float64
segment_osrm_distance  float64
segment_factor          float64
dtype: object

```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144867 entries, 0 to 144866
Data columns (total 24 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   data              144867 non-null   object 
 1   trip_creation_time 144867 non-null   object 
 2   route_schedule_uuid 144867 non-null   object 
 3   route_type         144867 non-null   object 
 4   trip_uuid          144867 non-null   object 
 5   source_center       144867 non-null   object 
 6   source_name         144574 non-null   object 
 7   destination_center 144867 non-null   object 
 8   destination_name    144606 non-null   object 
 9   od_start_time       144867 non-null   object 
 10  od_end_time         144867 non-null   object 
 11  start_scan_to_end_scan 144867 non-null   float64
 12  is_cutoff           144867 non-null   bool    
 13  cutoff_factor        144867 non-null   int64  
 14  cutoff_timestamp     144867 non-null   object 
 15  actual_distance_to_destination float64
 16  actual_time          144867 non-null   float64
 17  osrm_time            144867 non-null   float64
 18  osrm_distance         144867 non-null   float64
 19  factor               144867 non-null   float64
 20  segment_actual_time  144867 non-null   float64
 21  segment_osrm_time    144867 non-null   float64
 22  segment_osrm_distance 144867 non-null   float64
 23  segment_factor        144867 non-null   float64
dtypes: bool(1), float64(10), int64(1), object(12)
memory usage: 25.6+ MB

```

```

not_required_columns = ['is_cutoff', 'cutoff_factor', 'cutoff_timestamp', 'factor', 'segment_factor']
df=df.drop(columns = not_required_columns)

```

```
df.nunique()
```

data	2
trip_creation_time	14817
route_schedule_uuid	1504
route_type	2
trip_uuid	14817
source_center	1508
source_name	1498
destination_center	1481
destination_name	1468
od_start_time	26369
od_end_time	26369
start_scan_to_end_scan	1915
actual_distance_to_destination	144515
actual_time	3182
osrm_time	1531
osrm_distance	138046
segment_actual_time	747
segment_osrm_time	214
segment_osrm_distance	113799
dtype: int64	

```

df["data"] = df["data"].astype("category")
df["route_type"] = df["route_type"].astype("category")

df["trip_creation_time"] = pd.to_datetime(df["trip_creation_time"])

df["od_start_time"] = pd.to_datetime(df["od_start_time"])

df["od_end_time"] = pd.to_datetime(df["od_end_time"])

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144867 entries, 0 to 144866
Data columns (total 19 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   data             144867 non-null   category
 1   trip_creation_time 144867 non-null   datetime64[ns]
 2   route_schedule_uuid 144867 non-null   object  
 3   route_type        144867 non-null   category
 4   trip_uuid         144867 non-null   object  
 5   source_center     144867 non-null   object  
 6   source_name       144574 non-null   object  
 7   destination_center 144867 non-null   object  
 8   destination_name  144606 non-null   object  
 9   od_start_time    144867 non-null   datetime64[ns]
 10  od_end_time      144867 non-null   datetime64[ns]
 11  start_scan_to_end_scan 144867 non-null   float64
 12  actual_distance_to_destination 144867 non-null   float64
 13  actual_time       144867 non-null   float64
 14  osrm_time         144867 non-null   float64
 15  osrm_distance    144867 non-null   float64
 16  segment_actual_time 144867 non-null   float64
 17  segment_osrm_time 144867 non-null   float64
 18  segment_osrm_distance 144867 non-null   float64
dtypes: category(2), datetime64[ns](3), float64(8), object(6)
memory usage: 19.1+ MB

```

```

time_period = (df["od_end_time"].max(), df["trip_creation_time"].min())
time_period

(Timestamp('2018-10-08 03:00:24.353479'),
 Timestamp('2018-09-12 00:00:16.535741'))

```

Total 26 days of data are given in the dataset

## ▼ 1. Basic Data cleaning and exploration

```

import numpy as np
np.any(df.isnull())

```

True

## ▼ What is the number of null values present in each column?

```
df.isnull().sum()
```

data	0
trip_creation_time	0
route_schedule_uuid	0
route_type	0
trip_uuid	0
source_center	0
source_name	293
destination_center	0
destination_name	261
od_start_time	0
od_end_time	0
start_scan_to_end_scan	0
actual_distance_to_destination	0
actual_time	0

```

osrm_time          0
osrm_distance      0
segment_actual_time 0
segment_osrm_time   0
segment_osrm_distance 0
dtype: int64

```

```

df.duplicated()
len(df[df.duplicated()])

```

The number of duplicated values in the dataset are 0

```

missed_source_name = df.loc[df["source_name"].isnull(),"source_center"].unique()
missed_source_name

array(['IND342902A1B', 'IND577116AAA', 'IND282002AAD', 'IND465333A1B',
       'IND841301AAC', 'IND509103AAC', 'IND126116AAA', 'IND331022A1B',
       'IND505326AAB', 'IND852118A1B'], dtype=object)

missed_destination_name = df.loc[df["destination_name"].isnull(),"destination_center"].unique()
missed_destination_name

array(['IND342902A1B', 'IND577116AAA', 'IND282002AAD', 'IND465333A1B',
       'IND841301AAC', 'IND505326AAB', 'IND852118A1B', 'IND126116AAA',
       'IND509103AAC', 'IND221005A1A', 'IND250002AAC', 'IND331001A1C',
       'IND122015AAC'], dtype=object)

np.all(df.loc[df["source_name"].isnull()].isin(missed_destination_name))

False

```

## ✓ Treating missing names and destination names

```

count = 1
for i in missed_destination_name:
    df.loc[df['destination_center'] == i, 'destination_name'] = df.loc[df['destination_center'] == i, 'destination_name'].replace(
        count)
    count += 1
print(count)

14

d = {}
for i in missed_source_name:
    d[i] = df.loc[df['source_center'] == i, 'source_name'].unique()
for idx, val in d.items():
    if len(val) == 0:
        d[idx] = [f'location_{count}']
        count += 1
d2 = {}
for idx, val in d.items():
    d2[idx] = val[0]
for i, v in d2.items():
    print(i, v)

IND342902A1B location_1
IND577116AAA location_2
IND282002AAD location_3
IND465333A1B location_4
IND841301AAC location_5
IND509103AAC location_9
IND126116AAA location_8
IND331022A1B location_14
IND505326AAB location_6
IND852118A1B location_7

for i in missed_source_name:
    df.loc[df['source_center'] == i, 'source_name'] = df.loc[df['source_center'] == i, 'source_name'].replace(np.nan, d2[i])

df.isna().sum()

data              0
trip_creation_time 0
route_schedule_uuid 0

```

```

route_type          0
trip_uuid          0
source_center       0
source_name         0
destination_center  0
destination_name    0
od_start_time      0
od_end_time        0
start_scan_to_end_scan 0
actual_distance_to_destination 0
actual_time         0
osrm_time           0
osrm_distance       0
segment_actual_time 0
segment_osrm_time   0
segment_osrm_distance 0
dtype: int64

```

```
df.describe()
```

	start_scan_to_end_scan	actual_distance_to_destination	actual_time	osrm_time	osrm_distance	segment_actual_time
<b>count</b>	144867.000000	144867.000000	144867.000000	144867.000000	144867.000000	144867.000000
<b>mean</b>	961.262986	234.073372	416.927527	213.868272	284.771297	36.19611
<b>std</b>	1037.012769	344.990009	598.103621	308.011085	421.119294	53.57115
<b>min</b>	20.000000	9.000045	9.000000	6.000000	9.008200	-244.00000
<b>25%</b>	161.000000	23.355874	51.000000	27.000000	29.914700	20.00000
<b>50%</b>	449.000000	66.126571	132.000000	64.000000	78.525800	29.00000
<b>75%</b>	1634.000000	286.708875	513.000000	257.000000	343.193250	40.00000
<b>max</b>	7898.000000	1927.447705	4532.000000	1686.000000	2326.199100	3051.00000

```
df.describe(include=object)
```

	route_schedule_uuid	trip_uuid	source_center	source_name	destination_center	destination_name
<b>count</b>	144867	144867	144867	144867	144867	144867
<b>unique</b>	1504	14817	1508	1508	1481	1481
<b>top</b>	thanos::sroute:4029a8a2-6c74-4b7e-a6d8-f9e069f...	trip-153811219535896559	IND000000ACB	Gurgaon_Bilaspur_HB (Haryana)	IND000000ACB	Gurgaon_Bilaspur_HB (Haryana)

## ✓ merging of rows and aggregation of fields

```

df1 = df.groupby(by = ['trip_uuid', 'source_center', 'destination_center'], as_index = False).agg({'data' : 'first',
                                                 'route_type' : 'first',
                                                 'trip_creation_time' : 'first',
                                                 'source_name' : 'first',
                                                 'destination_name' : 'last',
                                                 'od_start_time' : 'first',
                                                 'od_end_time' : 'first',
                                                 'start_scan_to_end_scan' : 'first',
                                                 'actual_distance_to_destination' : 'last',
                                                 'actual_time' : 'last',
                                                 'osrm_time' : 'last',
                                                 'osrm_distance' : 'last',
                                                 'segment_actual_time' : 'sum',
                                                 'segment_osrm_time' : 'sum',
                                                 'segment_osrm_distance' : 'sum'})

```

```
df1.loc[df1["trip_uuid"] == "trip-153741093647649320"]
```

	trip_uuid	source_center	destination_center	data	route_type	trip_creation_time	source_name
10374	trip-153741093647649320	IND388121AAA	IND388620AAB	training	Carting	2018-09-20 02:35:36.476840	Anand_VUNagar_DC (Gujarat)
10375	trip-153741093647649320	IND388620AAB	IND388320AAA	training	Carting	2018-09-20 02:35:36.476840	Khambhat_MotvdDPP_D (Gujarat)

- Calculate the time taken between od\_start\_time and od\_end\_time and keep it as a feature. Drop the original columns, if required

```
df1['od_total_time'] = df1['od_end_time'] - df1['od_start_time']
df1.drop(columns = ['od_end_time', 'od_start_time'], inplace = True)
df1['od_total_time'] = df1['od_total_time'].apply(lambda x : round(x.total_seconds() / 60.0, 2))
df1['od_total_time'].head()
```

```
0    1260.60
1     999.51
2      58.83
3     122.78
4     834.64
Name: od_total_time, dtype: float64
```

```
df2 = df1.groupby(by="trip_uuid", as_index=False).agg({
    "source_center": "first",
    "destination_center": "last",
    "data": "first",
    "route_type": "first",
    "trip_creation_time": "first",
    "source_name": "first",
    "destination_name": "last",
    "od_total_time": "sum",
    "start_scan_to_end_scan": "sum",
    "actual_distance_to_destination": "sum",
    "actual_time": "sum",
    "osrm_time": "sum",
    "osrm_distance": "sum",
    "segment_actual_time": "sum",
    "segment_osrm_time": "sum",
    "segment_osrm_distance": "sum"})
```

```
df2.head()
```

	trip_uuid	source_center	destination_center	data	route_type	trip_creation_time	source_name
0	trip-153671041653548748	IND209304AAA	IND209304AAA	training	FTL	2018-09-12 00:00:16.535741	Kanpur_Central_H_6 (Uttar Pradesh)
1	trip-153671042288605164	IND561203AAB	IND561203AAB	training	Carting	2018-09-12 00:00:22.886430	Doddablpur_Chikadpp_D (Karnataka)
2	trip-153671043369099517	IND000000ACB	IND000000ACB	training	FTL	2018-09-12 00:00:33.691250	Gurgaon_Bilaspur_HB (Haryana)
3	trip-153671046011330457	IND400072AAB	IND401104AAA	training	Carting	2018-09-12 00:01:00.113710	Mumbai Hub (Maharashtra)
4	trip-153671052974046625	IND583101AAA	IND583119AAA	training	FTL	2018-09-12 00:02:09.740725	Bellary_Dc (Karnataka)

```
df2.loc[df2["trip_uuid"] == "trip-153741093647649320"]
```

	trip_uuid	source_center	destination_center	data	route_type	trip_creation_time	source_name	desti
5919	trip-153741093647649320	IND388121AAA	IND388320AAA	training	Carting	2018-09-20 02:35:36.476840	Anand_VUNagar_DC (Gujarat)	Anan

2. Build some features to prepare the data for actual analysis. Extract features from the below fields:

Source Name: Split and extract features out of destination. City-place-code (State)

```

def name_of_the_state(x):
    l = x.split("/")
    if len(l) == 1:
        return l[0]
    else:
        return l[1].replace("/", "")

df2["state"] = df2["source_name"].apply(name_of_the_state)
df2["state"].unique()

array(['Uttar Pradesh', 'Karnataka', 'Haryana', 'Maharashtra',
       'Tamil Nadu', 'Gujarat', 'Delhi', 'Telangana', 'Rajasthan',
       'Assam', 'Madhya Pradesh', 'West Bengal', 'Andhra Pradesh',
       'Punjab', 'Chandigarh', 'Goa', 'Jharkhand', 'Pondicherry',
       'Orissa', 'Uttarakhand', 'Himachal Pradesh', 'Kerala',
       'Arunachal Pradesh', 'Bihar', 'Chhattisgarh',
       'Dadra and Nagar Haveli', 'Jammu & Kashmir', 'Mizoram', 'Nagaland',
       'location_9', 'location_3', 'location_2', 'location_14',
       'location_7'], dtype=object)

def name_of_the_city(x):
    if 'location' in x:
        return 'unknown_city'
    else:
        l=x.split()[0].split("_")
        if 'CCU' in x:
            return 'Kolkata'
        elif 'MAA' in x.upper():
            return 'Chennai'
        elif ('HBR' in x.upper()) or ('BLR' in x.upper()):
            return 'Bengaluru'
        elif 'FBD' in x.upper():
            return 'Faridabad'
        elif 'BOM' in x.upper():
            return 'Mumbai'
        elif 'DEL' in x.upper():
            return 'Delhi'
        elif 'OK' in x.upper():
            return 'Delhi'
        elif 'GZB' in x.upper():
            return 'Ghaziabad'
        elif 'GGN' in x.upper():
            return 'Gurgaon'
        elif 'AMD' in x.upper():
            return 'Ahmedabad'
        elif 'CJB' in x.upper():
            return 'Coimbatore'
        elif 'HYD' in x.upper():
            return 'Hyderabad'
        return l[0]

df2["city"] = df2["source_name"].apply(name_of_the_city)
df2["city"].unique()[:50]

array(['Kanpur', 'Doddablpur', 'Gurgaon', 'Mumbai', 'Bellary', 'Chennai',
       'Bengaluru', 'Surat', 'Delhi', 'Pune', 'Faridabad', 'Shirala',
       'Hyderabad', 'Thirumalagiri', 'Gulbarga', 'Jaipur', 'Allahabad',
       'Guwahati', 'Narsinghpur', 'Shrirampur', 'Madakasira', 'Sonari',
       'Dindigul', 'Jalandhar', 'Chandigarh', 'Deoli', 'Pandharpur',
       'Kolkata', 'Bhandara', 'Kurnool', 'Bhiwandi', 'Bhatinda',
       'RoopNagar', 'Bantwal', 'Lalru', 'Kadi', 'Shahdol', 'Gangakher',
       'Durgapur', 'Vapi', 'Jamjodhpur', 'Jetpur', 'Mehsana', 'Jabalpur',
       'Junagadh', 'Gundlupet', 'Mysore', 'Goa', 'Bhopal', 'Sonipat'],
      dtype=object)

```

```

def name_of_the_place(x):
    if "location" in x:
        return x
    elif "HBR" in x:
        return "HBR Layout PC"
    else:
        l = x.split()[0].split("_", 1)
        if len(l) == 1:
            return "unknown_place"
        else:
            return l[1]

df2["place"] = df2["source_name"].apply(name_of_the_place)
df2["place"].unique()[:50]

array(['Central_H_6', 'ChikaDPP_D', 'Bilaspur_HB', 'unknown_place', 'Dc',
       'Poonamallee', 'Chrompet_DPC', 'HBR Layout PC', 'Central_D_12',
       'Lajpat_IP', 'North_D_3', 'Balabgarh_DPC', 'Central_DPP_3',
       'Shamshbd_H', 'Xroad_D', 'Nehrugnj_I', 'Central_I_7',
       'Central_H_1', 'Nangli_IP', 'North', 'KndliDPP_D', 'Central_D_9',
       'DavkhariD_D', 'Bandel_D', 'RTCStand_D', 'Central_DPP_1',
       'KGAirprt_HB', 'North_D_2', 'Central_D_1', 'DC', 'Mthurard_L',
       'Mullanpr_DC', 'Central_DPP_2', 'RajCmplx_D', 'Belaghata_DPC',
       'RjnaiDPP_D', 'AbbasNgr_I', 'Mankoli_HB', 'DPC', 'Airport_H',
       'Hub', 'Gateway_HB', 'Tathawde_H', 'ChotiHvl_DC', 'Trmltmp_D',
       'OnkarDPP_D', 'Mehmdpur_H', 'KaranNGR_D', 'Sohagpur_D',
       'Chrompet_L'], dtype=object)

```

#### ❖ Destination Name: Split and extract features out of destination. City-place-code (State)

```

df2["destination_state"] = df2["destination_name"].apply(name_of_the_state)
df2["destination_state"].unique()

array(['Uttar Pradesh', 'Karnataka', 'Haryana', 'Maharashtra',
       'Tamil Nadu', 'Gujarat', 'Delhi', 'Telangana', 'Rajasthan',
       'Madhya Pradesh', 'Assam', 'West Bengal', 'Andhra Pradesh',
       'Punjab', 'Chandigarh', 'Dadra and Nagar Haveli', 'Orissa',
       'Bihar', 'Jharkhand', 'Goa', 'Uttarakhand', 'Himachal Pradesh',
       'Kerala', 'Arunachal Pradesh', 'Mizoram', 'Chhattisgarh',
       'Jammu & Kashmir', 'Nagaland', 'Meghalaya', 'Tripura',
       'location_13', 'location_6', 'location_2', 'location_7',
       'location_3', 'location_5', 'location_12', 'location_11',
       'Daman & Diu'], dtype=object)

```

```

df2["destination_city"] = df2["destination_name"].apply(name_of_the_city)
df2["destination_city"].unique()[:50]

array(['Kanpur', 'Doddablpur', 'Gurgaon', 'Mumbai', 'Sandur', 'Chennai',
       'Bengaluru', 'Surat', 'Delhi', 'PNQ', 'Faridabad', 'Ratnagiri',
       'Bangalore', 'Hyderabad', 'Aland', 'Jaipur', 'Satna', 'Guwahati',
       'Bareli', 'Nashik', 'Hooghly', 'Sivasagar', 'Palani', 'Jalandhar',
       'Chandigarh', 'Yavatmal', 'Sangola', 'Kolkata', 'Savner',
       'Kurnool', 'Bhatinda', 'Bhiwandi', 'Barnala', 'Murbad', 'Kadaba',
       'Gulbarga', 'Naraingarh', 'Ludhiana', 'Kadi', 'Jabalpur',
       'Gangakher', 'Bankura', 'Silvassa', 'Porbandar', 'Jetpur',
       'Khammam', 'Mehsana', 'Katni', 'Una', 'Malavallili'], dtype=object)

```

```

df2["destination_place"] = df2["destination_name"].apply(name_of_the_place)
df2["destination_place"].unique()[:50]

array(['Central_H_6', 'ChikaDPP_D', 'Bilaspur_HB', 'MiraRd_IP',
       'WrdN1DPP_D', 'Poonamallee', 'Vandalur_Dc', 'HBR Layout PC',
       'Central_D_3', 'Bhogal', 'unknown_place', 'MjgaonRd_D',
       'Nelmngla_H', 'Uppal_I', 'RazaviRd_D', 'Central_I_7',
       'Central_I_2', 'Hub', 'SourvDPP_D', 'Varachha_DC', 'TgrniaRD_I',
       'DC', 'Gokulam_D', 'Babupaty_D', 'Bomsndra_HB', 'Alwal_I',
       'RjndraRd_D', 'Mehmdpur_H', 'Sanpada_I', 'JajuDPP_D',
       'Central_DPP_2', 'Dankuni_HB', 'Wagodha_D', 'AbbasNgr_I',
       'Balabgarh_DPC', 'DPC', 'Mankoli_HB', 'Shamshbd_H', 'SnkunDPP_D',
       'Kharar_DC', 'AnugrDPP_D', 'Nehrugnj_I', 'Ward2DPP_D',
       'MilrGanj_HB', 'KaranNGR_D', 'Adhartal_IP', 'Poonamallee_HB',
       'Busstand_D', 'BhowmDPP_D', 'Samrvrni_D'], dtype=object)

```

#### ❖ Trip\_creation\_time: Extract features like month, year and day etc

```

df2['trip_creation_hour'] = df2['trip_creation_time'].dt.hour
df2['trip_creation_hour'].head()

0    0
1    0
2    0
3    0
4    0
Name: trip_creation_hour, dtype: int64

df2['trip_creation_date'] = pd.to_datetime(df2['trip_creation_time'].dt.date)
df2['trip_creation_date'].head()

0   2018-09-12
1   2018-09-12
2   2018-09-12
3   2018-09-12
4   2018-09-12
Name: trip_creation_date, dtype: datetime64[ns]

df2['trip_creation_day'] = df2['trip_creation_time'].dt.day
df2['trip_creation_day'].head()

0    12
1    12
2    12
3    12
4    12
Name: trip_creation_day, dtype: int64

df2['trip_creation_week'] = df2['trip_creation_time'].dt.isocalendar().week
df2['trip_creation_week'].head()

0    37
1    37
2    37
3    37
4    37
Name: trip_creation_week, dtype: UInt32

df2['trip_creation_month'] = df2['trip_creation_time'].dt.month
df2['trip_creation_month'].head()

0    9
1    9
2    9
3    9
4    9
Name: trip_creation_month, dtype: int64

df2['trip_creation_year'] = df2['trip_creation_time'].dt.year
df2['trip_creation_year'].head()

0    2018
1    2018
2    2018
3    2018
4    2018
Name: trip_creation_year, dtype: int64

```

## ✓ Structure of the Data after Data Cleaning

```
df2.shape
```

```
(14817, 29)
```

```
df2.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14817 entries, 0 to 14816
Data columns (total 29 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   trip_uuid        14817 non-null   object 
 1   source_center    14817 non-null   object 
 2   destination_center 14817 non-null   object 

```

```

3   data          14817 non-null category
4   route_type    14817 non-null category
5   trip_creation_time 14817 non-null datetime64[ns]
6   source_name    14817 non-null object
7   destination_name 14817 non-null object
8   od_total_time  14817 non-null float64
9   start_scan_to_end_scan 14817 non-null float64
10  actual_distance_to_destination 14817 non-null float64
11  actual_time     14817 non-null float64
12  osrm_time       14817 non-null float64
13  osrm_distance   14817 non-null float64
14  segment_actual_time 14817 non-null float64
15  segment_osrm_time 14817 non-null float64
16  segment_osrm_distance 14817 non-null float64
17  state           14817 non-null object
18  city            14817 non-null object
19  place           14817 non-null object
20  destination_state 14817 non-null object
21  destination_city 14817 non-null object
22  destination_place 14817 non-null object
23  trip_creation_hour 14817 non-null int64
24  trip_creation_date 14817 non-null datetime64[ns]
25  trip_creation_day 14817 non-null int64
26  trip_creation_week 14817 non-null UInt32
27  trip_creation_month 14817 non-null int64
28  trip_creation_year 14817 non-null int64
dtypes: UInt32(1), category(2), datetime64[ns](2), float64(9), int64(4), object(11)
memory usage: 3.0+ MB

```

df2.describe().T

	count	mean	std	min	25%	50%	75%	max
<b>od_total_time</b>	14817.0	531.69763	658.868223	23.46	149.93	280.77	638.2	7898.55
<b>start_scan_to_end_scan</b>	14817.0	530.810016	658.705957	23.0	149.0	280.0	637.0	7898.0
<b>actual_distance_to_destination</b>	14817.0	164.477838	305.388147	9.002461	22.837239	48.474072	164.583208	2186.531787
<b>actual_time</b>	14817.0	357.143754	561.396157	9.0	67.0	149.0	370.0	6265.0
<b>osrm_time</b>	14817.0	161.384018	271.360995	6.0	29.0	60.0	168.0	2032.0
<b>osrm_distance</b>	14817.0	204.344689	370.395573	9.0729	30.8192	65.6188	208.475	2840.081
<b>segment_actual_time</b>	14817.0	353.892286	556.247965	9.0	66.0	147.0	367.0	6230.0
<b>segment_osrm_time</b>	14817.0	180.949787	314.542047	6.0	31.0	65.0	185.0	2564.0
<b>segment_osrm_distance</b>	14817.0	223.201161	416.628374	9.0729	32.6545	70.1544	218.8024	3523.6324
<b>trip_creation_hour</b>	14817.0	12.449821	7.986553	0.0	4.0	14.0	20.0	23.0
<b>trip_creation_day</b>	14817.0	18.37079	7.893275	1.0	14.0	19.0	25.0	30.0
<b>trip_creation_week</b>	14817.0	38.295944	0.967872	37.0	38.0	38.0	39.0	40.0
<b>trip_creation_month</b>	14817.0	9.120672	0.325757	9.0	9.0	9.0	9.0	10.0
<b>trip_creation_year</b>	14817.0	2018.0	0.0	2018.0	2018.0	2018.0	2018.0	2018.0

df2.describe(include=object).T


	count	unique	top	freq
<b>trip_uuid</b>	14817	14817	trip-153671041653548748	1
<b>source_center</b>	14817	938	IND000000ACB	1063
<b>destination_center</b>	14817	1042	IND000000ACB	821
<b>source_name</b>	14817	938	Gurgaon_Bilaspur_HB (Haryana)	1063
<b>destination_name</b>	14817	1042	Gurgaon_Bilaspur_HB (Haryana)	821
<b>state</b>	14817	34	Maharashtra	2714
<b>city</b>	14817	690	Mumbai	1442
<b>place</b>	14817	761	Bilaspur_HB	1063
<b>destination_state</b>	14817	39	Maharashtra	2561
<b>destination_city</b>	14817	806	Mumbai	1548
<b>destination_place</b>	14817	850	Bilaspur_HB	821

## ✓ Trips created on hourly basis

```
df_hourly_trip = df2.groupby(by="trip_creation_hour")["trip_uuid"].count().to_frame().reset_index()
df_hourly_trip.head()
```

	trip_creation_hour	trip_uuid
0	0	994
1	1	750
2	2	702
3	3	652
4	4	636

```
import matplotlib.pyplot as plt
import seaborn as sns
sns.lineplot(data=df_hourly_trip,
              x=df_hourly_trip["trip_creation_hour"],
              y=df_hourly_trip["trip_uuid"])
plt.yticks(np.arange(0,1400,200))
plt.xticks(np.arange(0,26,2))
plt.xlabel("hour")
plt.ylabel("count of trips")
plt.grid("both")
```



## ✓ daily trips


```
df_daily_trip = df2.groupby(by="trip_creation_day")["trip_uuid"].count().to_frame().reset_index()
df_daily_trip.head()
```

	trip_creation_day	trip_uuid
0	1	605
1	2	552
2	3	631
3	12	747
4	13	750

```

plt.figure(figsize=(10,6))
sns.lineplot(data=df_daily_trip,
              x=df_daily_trip["trip_creation_day"],
              y=df_daily_trip["trip_uuid"])
plt.xlabel("Day")
plt.xticks(np.arange(0,32))
plt.ylabel("count of trips")
plt.grid("both")

```



## ❖ Weekly trips

```

df_weekly_trip = df2.groupby(by="trip_creation_week")["trip_uuid"].count().to_frame().reset_index()
df_weekly_trip.head()


```

	trip_creation_week	trip_uuid
0	37	3608
1	38	5004
2	39	4417
3	40	1788

```

sns.lineplot(data=df_weekly_trip,
              x=df_weekly_trip["trip_creation_week"],
              y=df_weekly_trip["trip_uuid"])
plt.xlabel("week")
plt.ylabel("count of trips")
plt.grid("both")

```




## ▼ Monthly trips

```
df_monthly_trip = df2.groupby(by="trip_creation_month")["trip_uuid"].count().to_frame().reset_index()
df_monthly_trip["percentage"] = np.round(df_monthly_trip["trip_uuid"]*100/df_monthly_trip["trip_uuid"].sum(),2)
df_monthly_trip
```

	trip_creation_month	trip_uuid	percentage
0	9	13029	87.93
1	10	1788	12.07

```
plt.pie(x=df_monthly_trip["percentage"],
        labels=["Sep","Oct"],
        autopct = '%.2f%')
plt.show()
```




## ▼ Distribution of orders (Route wise)

```
df_routewise_trip=df2.groupby(by="route_type")["trip_uuid"].count().to_frame().reset_index()
df_routewise_trip["percentage"] = np.round(df_routewise_trip["trip_uuid"]*100/df_routewise_trip["trip_uuid"].sum(),2)
df_routewise_trip
```

	route_type	trip_uuid	percentage
0	Carting	8908	60.12
1	FTL	5909	39.88

```
plt.pie(x=df_routewise_trip["percentage"],
        labels=["Carting","FTL"],
        autopct = "%.2f%")
plt.show()
```




## ▼ Top 20 States with Maximum Trips

```
df_source_state = df2.groupby(by="state")["trip_uuid"].count().to_frame().reset_index()
df_source_state["percentage"] = np.round(df_source_state["trip_uuid"]*100/df_source_state["trip_uuid"].sum(),2)
df_source_state = df_source_state.sort_values(by="trip_uuid",ascending=False)[:20]
df_source_state.head()
```

	state	trip_uuid	percentage
17	Maharashtra	2714	18.32
14	Karnataka	2143	14.46
10	Haryana	1838	12.40
24	Tamil Nadu	1039	7.01
25	Telangana	781	5.27

```
plt.figure(figsize=(10,8))
sns.barplot(data = df_source_state,
            x = df_source_state["trip_uuid"],
            y = df_source_state["state"])
plt.plot()
```

[]




#### ▼ Top 20 cities with Maximum trips

```
df_source_city = df2.groupby(by="city")["trip_uuid"].count().to_frame().reset_index()
df_source_city["percentage"] = np.round(df_source_city["trip_uuid"]*100/df_source_city["trip_uuid"].sum(),2)
df_source_city = df_source_city.sort_values(by="trip_uuid",ascending=False)[:20]
df_source_city.head()
```

	city	trip_uuid	percentage
439	Mumbai	1442	9.73
237	Gurgaon	1165	7.86
169	Delhi	883	5.96
79	Bengaluru	726	4.90
100	Bhiwandi	697	4.70

```
plt.figure(figsize=(10,8))
sns.barplot(data = df_source_city,
            x = df_source_city["trip_uuid"],
            y = df_source_city["city"])
plt.plot()
```


[]



```

numerical_columns = ['od_total_time', 'start_scan_to_end_scan', 'actual_distance_to_destination',
                     'actual_time', 'osrm_time', 'osrm_distance', 'segment_actual_time',
                     'segment_osrm_time', 'segment_osrm_distance']
sns.pairplot(data = df2,
              vars = numerical_columns,
              kind = 'reg',
              hue = 'route_type',
              markers = '.')
plt.plot()

```




```
df_corr = df2[numerical_columns].corr()
df_corr
```

	od_total_time	start_scan_to_end_scan	actual_distance_to_destination	actual_time	osrm_time	osi
od_total_time	1.000000	0.999999	0.918222	0.961094	0.926516	
start_scan_to_end_scan	0.999999	1.000000	0.918308	0.961147	0.926571	
actual_distance_to_destination	0.918222	0.918308	1.000000	0.953757	0.993561	
actual_time	0.961094	0.961147	0.953757	1.000000	0.958593	
osrm_time	0.926516	0.926571	0.993561	0.958593	1.000000	
osrm_distance	0.924219	0.924299	0.997264	0.959214	0.997580	
segment_actual_time	0.961119	0.961171	0.952821	0.999989	0.957765	
segment_osrm_time	0.918490	0.918561	0.987538	0.953872	0.993259	
segment_osrm_distance	0.919199	0.919291	0.993061	0.956967	0.991608	

```
plt.figure(figsize = (15, 10))
sns.heatmap(data = df_corr, vmin = -1, vmax = 1, annot = True)
plt.plot()
```

[]



### 3. In-depth analysis and feature engineering:

- Compare the difference between od\_total\_time and start\_scan\_to\_end\_scan. Do hypothesis testing/ Visual analysis to check.

# Steps of Feature Engineering:

```

# Step 1: Set up Null & Alternate Hypothesis
# H0(Null Hypothesis): mean of od_total_time(Total trip time) and start_scan_to_end (Expected total trip time) are same.
# HA(Alternate Hypothesis): mean of od_total_time(Total trip time) and start_scan_to_end (Expected total trip time) are different

# step 2: Basic assumptions for Hypothesis
# Distribution using QQ plot
# Homogeneity of variances using lavene's test

# step 3: Define Test Statistics
# If the assumptions of T Test are met then we can proceed performing T Test for independent samples else
# we will perform the non parametric test equivalent to T Test for independent sample
# i.e., Mann-Whitney U rank test for two independent samples.

# step 4: compute the p-value and fix value of alpha
# alpha = 0.05

# step 5: Based on p-value we will reject or accept H0.
# p-val > alpha : Accept H0
# p-val < alpha : Reject H0

# checking mean, median and mode are equal
df2[["od_total_time","start_scan_to_end_scan"]].describe()


```

	od_total_time	start_scan_to_end_scan
count	14817.000000	14817.000000
mean	531.697630	530.810016
std	658.868223	658.705957
min	23.460000	23.000000
25%	149.930000	149.000000
50%	280.770000	280.000000
75%	638.200000	637.000000
max	7898.550000	7898.000000

## ❖ Histogram

```
# If it is following normal distribution, it will form a shape like bell.
```

```
sns.histplot(df2['od_total_time'], element = 'step', color = 'green')
sns.histplot(df2['start_scan_to_end_scan'], element = 'step', color = 'pink')
plt.legend(['od_total_time', 'start_scan_to_end_scan'])
plt.plot()
```




## ❖ Distribution check using QQ Plot

```
import scipy.stats as spy
plt.figure(figsize = (15, 6))
plt.subplot(1, 2, 1)
plt.suptitle('QQ plots for od_total_time and start_scan_to_end_scan')
spy.probplot(df2['od_total_time'], plot = plt, dist = 'norm')
plt.title('QQ plot for od_total_time')
plt.subplot(1, 2, 2)
spy.probplot(df2['start_scan_to_end_scan'], plot = plt, dist = 'norm')
plt.title('QQ plot for start_scan_to_end_scan')
plt.plot()
```

[]

QQ plots for od\_total\_time and start\_scan\_to\_end\_scan



From the above both plots, we can infer that it will not follow normal distribution

#### ✓ Applying shapiro-wilk test for normality

$H_0$  : The sample follows normal distribution

$H_1$  : The sample does not follow normal distribution

alpha = 0.05

```
test_stat, p_value = spy.shapiro(df2['od_total_time'].sample(5000))
print('p-value', p_value)
if p_value < 0.05:
    print('The sample does not follow normal distribution')
else:
    print('The sample follows normal distribution')

p-value 0.0
The sample does not follow normal distribution

test_stat, p_value = spy.shapiro(df2['start_scan_to_end_scan'].sample(5000))
print('p-value', p_value)
if p_value < 0.05:
    print('The sample does not follow normal distribution')
else:
    print('The sample follows normal distribution')

p-value 0.0
The sample does not follow normal distribution
```

#### ✓ Using box-cox transformation

```
transformed_start_scan_to_end_scan = spy.boxcox(df2['start_scan_to_end_scan'])[0]
test_stat, p_value = spy.shapiro(transformed_start_scan_to_end_scan)
print('p-value', p_value)
if p_value < 0.05:
    print('The sample does not follow normal distribution')
else:
    print('The sample follows normal distribution')

p-value 1.0471322892609475e-24
The sample does not follow normal distribution
/usr/local/lib/python3.10/dist-packages/scipy/stats/_morestats.py:1882: UserWarning: p-value may not be accurate for N > 5000
warnings.warn("p-value may not be accurate for N > 5000.")
```

## ✓ Homogeneity of variances using lavene's test

```
# Null Hypothesis(H0) - Homogenous Variance
# Alternate Hypothesis(HA) - Non Homogenous Variance

test_stat, p_value = spy.levene(df2['od_total_time'], df2['start_scan_to_end_scan'])
print('p-value', p_value)
if p_value < 0.05:
    print('The samples do not have Homogenous Variance')
else:
    print('The samples have Homogenous Variance')

p-value 0.9668007217581142
The samples have Homogenous Variance
```

- ✓ We can not perform T-test here, because the samples are not normally distributed here, we can perform non-parametric equivalent test like Mann-Whitney U rank here for two-independent samples

```
# H0: mean of od_total_time and start_scan_to_end_scan are same.
# HA: mean of od_total_time and start_scan_to_end_scan are different
alpha = 0.05
test_stat, p_value = spy.mannwhitneyu(df2['od_total_time'], df2['start_scan_to_end_scan'])
print('P-value :',p_value)
if p_value > 0.05:
    print("mean of od_total_time and start_scan_to_end_scan are same")
else:
    print("mean of od_total_time and start_scan_to_end_scan are different")


P-value : 0.7815123224221716
mean of od_total_time and start_scan_to_end_scan are same
```

mean of time taken by od\_total\_time and start\_scan\_to\_end\_scan are same

## ✓ Hypothesis testing / visual analysis between actual\_time and osrm\_time

### ✓ Visual tests to know if the sample follows normal distribution

```
sns.histplot(df2["actual_time"], color="red")
sns.histplot(df2["osrm_time"], color="yellow")
plt.legend(["actual_time", "osrm_time"])
plt.show()
```




## ✓ QQ Plot

```
plt.figure(figsize = (15, 6))
plt.subplot(1, 2, 1)
plt.suptitle('QQ plots for actual_time and osrm_time')
spy.probplot(df2['actual_time'], plot = plt, dist = 'norm')
plt.title('QQ plot for actual_time')
plt.subplot(1, 2, 2)
spy.probplot(df2['osrm_time'], plot = plt, dist = 'norm')
plt.title('QQ plot for osrm_time')
plt.plot()
```

[]

QQ plots for actual\_time and osrm\_time



It can be inferred that samples do not follow normal distribution

## ✓ Applying shapiro-wilk test for normality

```
test_stat, p_value = spy.shapiro(df2['actual_time'].sample(5000))
print('p-value', p_value)
if p_value < 0.05:
    print('The sample does not follow normal distribution')
else:
    print('The sample follows normal distribution')

p-value 0.0
The sample does not follow normal distribution

test_stat, p_value = spy.shapiro(df2['osrm_time'].sample(5000))
print('p-value', p_value)
if p_value < 0.05:
    print('The sample does not follow normal distribution')
else:
    print('The sample follows normal distribution')

p-value 0.0
The sample does not follow normal distribution
```

## ✓ boxcox Transformation

```

transformed_actual_time = spy.boxcox(df2['actual_time'])[0]
test_stat, p_value = spy.shapiro(transformed_actual_time)
print('p-value', p_value)
if p_value < 0.05:
    print('The sample does not follow normal distribution')
else:
    print('The sample follows normal distribution')

p-value 1.021792743086169e-28
The sample does not follow normal distribution
/usr/local/lib/python3.10/dist-packages/scipy/stats/_morestats.py:1882: UserWarning: p-value may not be accurate for N > 5000
    warnings.warn("p-value may not be accurate for N > 5000.")

```

```

transformed_osrm_time = spy.boxcox(df2['osrm_time'])[0]
test_stat, p_value = spy.shapiro(transformed_osrm_time)
print('p-value', p_value)
if p_value < 0.05:
    print('The sample does not follow normal distribution')
else:
    print('The sample follows normal distribution')

p-value 3.543600614978861e-35
The sample does not follow normal distribution

```

- Even after applying the boxcox transformation on each of the "actual\_time" and "osrm\_time" columns, the distributions do not follow normal distribution.
- Homogeneity of Variances using **Lavene's test**

```

# Null Hypothesis(H0) - Homogenous Variance

# Alternate Hypothesis(HA) - Non Homogenous Variance

test_stat, p_value = spy.levene(df2['actual_time'], df2['osrm_time'])
print('p-value', p_value)
if p_value < 0.05:
    print('The samples do not have Homogenous Variance')
else:
    print('The samples have Homogenous Variance')

p-value 1.871297993683208e-220
The samples do not have Homogenous Variance

```

- Since the samples do not follow any of the assumptions T-Test cannot be applied here, we can perform its non parametric equivalent test i.e., Mann-Whitney U rank test for two independent samples.

```

test_stat, p_value = spy.mannwhitneyu(df2['actual_time'], df2['osrm_time'])
print('p-value', p_value)
if p_value < 0.05:
    print('mean of actual_time and osrm_time are not similar')
else:
    print('mean of actual_time and osrm_time are similar')

p-value 0.0
mean of actual_time and osrm_time are not similar

```

mean of actual\_time and osrm\_time are not similar


- Do Hypothesis testing between actual\_time and segment\_actual\_time
- Histogram

```

sns.histplot(df2['actual_time'], element = 'step', color = 'blue')
sns.histplot(df2['segment_actual_time'], element = 'step', color = 'red')
plt.legend(['actual_time', 'segment_actual_time'])
plt.plot()

```

[]



- Distribution check using **QQ Plot**


```

plt.figure(figsize = (12, 6))
plt.subplot(1, 2, 1)
plt.suptitle('QQ plots for actual_time and segment_actual_time')
sns.probplot(df2['actual_time'], plot = plt, dist = 'norm')
plt.title('QQ plot for actual_time')
plt.subplot(1, 2, 2)
sns.probplot(df2['segment_actual_time'], plot = plt, dist = 'norm')
plt.title('QQ plot for segment_actual_time')
plt.plot()

```

[]

QQ plots for actual\_time and segment\_actual\_time



- ✓ Samples do not follow normal distribution.

- Applying Shapiro-Wilk test for normality

$H_0$  : The sample **follows normal distribution**

$H_1$  : The sample **does not follow normal distribution**

alpha = 0.05

```
test_stat, p_value = spy.shapiro(df2['actual_time'].sample(5000))
print('p-value', p_value)
if p_value < 0.05:
    print('The sample does not follow normal distribution')
else:
    print('The sample follows normal distribution')

p-value 0.0
The sample does not follow normal distribution

test_stat, p_value = spy.shapiro(df2['segment_actual_time'].sample(5000))
print('p-value', p_value)
if p_value < 0.05:
    print('The sample does not follow normal distribution')
else:
    print('The sample follows normal distribution')

p-value 0.0
The sample does not follow normal distribution
```

- Transforming the data using boxcox transformation to check if the transformed data follows normal distribution.

```
transformed_actual_time = spy.boxcox(df2['actual_time'])[0]
test_stat, p_value = spy.shapiro(transformed_actual_time)
print('p-value', p_value)
if p_value < 0.05:
    print('The sample does not follow normal distribution')
else:
    print('The sample follows normal distribution')

p-value 1.021792743086169e-28
The sample does not follow normal distribution
/usr/local/lib/python3.10/dist-packages/scipy/stats/_morestats.py:1882: UserWarning: p-value may not be accurate for N > 5000
    warnings.warn("p-value may not be accurate for N > 5000.")
```

```
transformed_segment_actual_time = spy.boxcox(df2['segment_actual_time'])[0]
test_stat, p_value = spy.shapiro(transformed_segment_actual_time)
print('p-value', p_value)
if p_value < 0.05:
    print('The sample does not follow normal distribution')
else:
    print('The sample follows normal distribution')

p-value 5.696120172016859e-29
The sample does not follow normal distribution
```

- Even after applying the boxcox transformation on each of the "actual\_time" and "segment\_actual\_time" columns, the distributions do not follow normal distribution.

Homogeneity of Variances using Lavene's test

```
# Null Hypothesis(H0) - Homogenous Variance

# Alternate Hypothesis(HA) - Non Homogenous Variance

test_stat, p_value = spy.levene(df2['actual_time'], df2['segment_actual_time'])
print('p-value', p_value)

if p_value < 0.05:
    print('The samples do not have Homogenous Variance')
else:
    print('The samples have Homogenous Variance ')
```

```
p-value 0.6955022668700895
```

The samples have Homogenous Variance

- Since the samples do not come from normal distribution T-Test cannot be applied here, we can perform its non parametric equivalent test i.e., Mann-Whitney U rank test for two independent samples.

```
test_stat, p_value = spy.mannwhitneyu(df2['actual_time'], df2['segment_actual_time'])
print('p-value', p_value)
if p_value < 0.05:
    print('mean of actual_time and segment_actual_time are not similar')
else:
    print('mean of actual_time and segment_actual_time are similar')

p-value 0.4164235159622476
mean of actual_time and segment_actual_time are similar
```

mean of actual\_time and segment\_actual\_time are similar

- Find outliers in the numerical variables (you might find outliers in almost all the variables), and check it using visual analysis

```
numerical_columns = ['od_total_time', 'start_scan_to_end_scan', 'actual_distance_to_destination',
                     'actual_time', 'osrm_time', 'osrm_distance', 'segment_actual_time',
                     'segment_osrm_time', 'segment_osrm_distance']
df2[numerical_columns].describe().T
```


	count	mean	std	min	25%	50%	75%	max
<b>od_total_time</b>	14817.0	531.697630	658.868223	23.460000	149.930000	280.770000	638.200000	7898.550000
<b>start_scan_to_end_scan</b>	14817.0	530.810016	658.705957	23.000000	149.000000	280.000000	637.000000	7898.000000
<b>actual_distance_to_destination</b>	14817.0	164.477838	305.388147	9.002461	22.837239	48.474072	164.583208	2186.531787
<b>actual_time</b>	14817.0	357.143754	561.396157	9.000000	67.000000	149.000000	370.000000	6265.000000
<b>osrm_time</b>	14817.0	161.384018	271.360995	6.000000	29.000000	60.000000	168.000000	2032.000000
<b>osrm_distance</b>	14817.0	204.344689	370.395573	9.072900	30.819200	65.618800	208.475000	2840.081000
<b>segment_actual_time</b>	14817.0	353.892286	556.247965	9.000000	66.000000	147.000000	367.000000	6230.000000
<b>segment_osrm_time</b>	14817.0	180.949787	314.542047	6.000000	31.000000	65.000000	185.000000	2564.000000
<b>segment_osrm_distance</b>	14817.0	223.201161	416.628374	9.072900	32.654500	70.154400	218.802400	3523.632400

```
import matplotlib as mpl
colors_list = list(mpl.colors.cnames)

clr = np.random.choice(colors_list)
clr


'lightslategrey'

plt.figure(figsize = (18, 15))
for i in range(len(numerical_columns)):
    plt.subplot(3, 3, i + 1)
    clr = np.random.choice(list(mpl.colors.cnames))
    sns.histplot(df2[numerical_columns[i]], bins = 1000, kde = True, color = clr)
    plt.title(f"Distribution of {numerical_columns[i]} column")
    plt.plot()
```



- As we can see from the above plots, all the histograms of numerical columns are right skewed, that means there are outliers present in the columns that needs to be cleaned.

```
plt.figure(figsize = (18, 15))
for i in range(len(numerical_columns)):
    plt.subplot(3, 3, i + 1)
    clr = np.random.choice(list(mpl.colors.cnames))
    sns.boxplot(df2[numerical_columns[i]], color = clr)
    plt.title(f"Distribution of {numerical_columns[i]} column")
    plt.plot()
```



It can be clearly seen in the above plots that there are outliers in all the numerical columns that can be treated.

```
# Detecting Outliers

for i in numerical_columns:
    Q1 = np.quantile(df2[i], 0.25)
    Q3 = np.quantile(df2[i], 0.75)
    IQR = Q3 - Q1
    LB = Q1 - 1.5 * IQR
    UB = Q3 + 1.5 * IQR
    outliers = df2.loc[(df2[i] < LB) | (df2[i] > UB)]
    print('Column :', i)
    print(f'Q1 : {Q1}')
    print(f'Q3 : {Q3}')
    print(f'IQR : {IQR}')
    print(f'LB : {LB}')
    print(f'UB : {UB}')
    print(f'Number of outliers : {outliers.shape[0]}')
    print('-----')
```

Column : od\_total\_time

Q1 : 149.93

Q3 : 638.2

IQR : 488.2700000000004

LB : -582.4750000000001

UB : 1370.605

Number of outliers : 1266

-----

Column : start\_scan\_to\_end\_scan

Q1 : 149.0

Q3 : 637.0

IQR : 488.0

```

LB : -583.0
UB : 1369.0
Number of outliers : 1267
-----
Column : actual_distance_to_destination
Q1 : 22.83723905859321
Q3 : 164.58320763841138
IQR : 141.74596857981817
LB : -189.78171381113404
UB : 377.2021685081386
Number of outliers : 1449
-----
Column : actual_time
Q1 : 67.0
Q3 : 370.0
IQR : 303.0
LB : -387.5
UB : 824.5
Number of outliers : 1643
-----
Column : osrm_time
Q1 : 29.0
Q3 : 168.0
IQR : 139.0
LB : -179.5
UB : 376.5
Number of outliers : 1517
-----
Column : osrm_distance
Q1 : 30.8192
Q3 : 208.475
IQR : 177.6558
LB : -235.6645
UB : 474.9587
Number of outliers : 1524
-----
Column : segment_actual_time
Q1 : 66.0
Q3 : 367.0
IQR : 301.0
LB : -385.5
UB : 818.5
Number of outliers : 1643
-----
Column : segment_osrm_time

```

One Hot encoding is not right suitable in route\_type or data column, that is why we used label encoding. Label encoding also has some disadvantages like it expands our dataset, which is not needed.

#### ✓ Label Encoding of categorical variables

```

# Get value counts before Label encoding

df2['route_type'].value_counts()

Carting    8908
FTL        5909
Name: route_type, dtype: int64

# Perform label encoding on categorical column route type

from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
df2['route_type'] = label_encoder.fit_transform(df2['route_type'])

# Get value counts after Label encoding

df2['route_type'].value_counts()

0    8908
1    5909
Name: route_type, dtype: int64

# Get value counts of categorical variable 'data' before Label encoding

df2['data'].value_counts()

```

```

training    10654
test        4163
Name: data, dtype: int64

label_encoder = LabelEncoder()
df2['data'] = label_encoder.fit_transform(df2['data'])

# Get value counts after label encoding

df2['data'].value_counts()

1    10654
0     4163
Name: data, dtype: int64

```

- Normalize/ Standardize the numerical features using MinMaxScaler or StandardScaler.

```
from sklearn.preprocessing import MinMaxScaler
```

```

plt.figure(figsize = (10, 6))
scaler = MinMaxScaler()
scaled = scaler.fit_transform(df2['od_total_time'].to_numpy().reshape(-1, 1))
sns.histplot(scaled)
plt.title(f"{scaled}")
plt.legend('od_total_time')
plt.plot()


```

[]

```

[[0.2840158 ]
[0.02008231]
[0.49661655]
...
[0.05062291]
[0.04127699]
[0.04202365]]

```




```

plt.figure(figsize = (10, 6))
scaler = MinMaxScaler()
scaled = scaler.fit_transform(df2['start_scan_to_end_scan'].to_numpy().reshape(-1, 1))
sns.histplot(scaled)
plt.title(f"{scaled}")
plt.plot()

```

[]


```
[[0.28393651]
 [0.01993651]
 [0.49650794]
 ...
 [0.05053968]
 [0.04114286]
 [0.04190476]]
```



```
plt.figure(figsize = (10, 6))
scaler = MinMaxScaler()
scaled = scaler.fit_transform(df2['actual_distance_to_destination'].to_numpy().reshape(-1, 1))
sns.histplot(scaled)
plt.title(f"{scaled}")
plt.plot()
```

[]


```
[[0.37461282]
 [0.02947581]
 [0.8809993 ]
 ...
 [0.01363122]
 [0.05773579]
 [0.02621277]]
```



```
plt.figure(figsize = (10, 6))
scaler = MinMaxScaler()
scaled = scaler.fit_transform(df2['segment_osrm_time'].to_numpy().reshape(-1, 1))
sns.histplot(scaled)
plt.title(f"{scaled}")
plt.plot()
```

[]

```
[[0.39171228]
 [0.02306489]
 [0.75645035]
 ...
 [0.03205629]
 [0.08405004]
 [0.02384676]]
```



```
plt.figure(figsize = (10, 6))
scaler = MinMaxScaler()
scaled = scaler.fit_transform(df2['segment_osrm_distance'].to_numpy().reshape(-1, 1))
sns.histplot(scaled)
plt.title(f"{df2['segment_osrm_distance']}")
```

[]

```

0      1320.4733
1      84.1894
2      2545.2678
3      19.8766
4      146.7919
...


```

```

14812    64.8551
14813    16.0883
14814    104.8866
14815    223.5324
14816    80.5787

```

Name: segment\_osrm\_distance, Length: 14817, dtype: float64



## ▼ Standard Scaler

```

from sklearn.preprocessing import StandardScaler

plt.figure(figsize = (10, 6))
scaler = StandardScaler()
scaled = scaler.fit_transform(df2['osrm_time'].to_numpy().reshape(-1, 1))
sns.histplot(scaled)
plt.title(f"{scaled}")
plt.plot()

```

[ ]

```
[[ 2.04758511]
 [-0.34414367]
 [ 5.81759828]
 ...
 [-0.41784872]
 [ 0.06491934]
 [-0.34414367]]
```



```
plt.figure(figsize = (10, 6))
scaler = StandardScaler()
scaled = scaler.fit_transform(df2['osrm_distance'].to_numpy().reshape(-1, 1))
sns.histplot(scaled)
plt.title(f"{scaled}")
plt.plot()
```